

CS422 Project: Searching for Markovian Malware C&C on Twitter

Maxime Augier and Gowthami Ramasamy

May 7, 2014

0.1 Objectives

The main objective was to assess how realistic is it to use a Markov language model to build a stealthy control channel for malware, piggybacking on the Twitter service, and as a side effect, to study other anomalies of Twitter traffic.

A botnet is typically employed to launch Denial-of-Service attacks on popular websites, then extorting protection money from the site owners. In its simplest form, a botnet command would consist of a target website, maybe a few control flags (3-5 bits) and an attack duration (to get a 5 minute granularity over 1 week, 11 bits are sufficient). According to the list of the 500 most popular domains available on <http://moz.com/top500/domains/csv>, the median length of a domain name is 14 characters and the 95th percentile is at 20. A domain name consists of the characters a-z, 0-9, dash and dot. With character-wise encoding, and assuming characters in domain names are evenly distributed (quite a pessimistic assumption) we need 6 bits to per domain character, so 120 bits for a domain name, 136 after adding some control bits.

We thus consider a control channel “realistic” if it can fit 136 bits in a single message.

0.2 Project setup

All the project code is available on github at <http://github.com/maugier/cs422.git>.

0.3 Data collection

As a starting point, we used a sample of the twitter stream archive provided on <http://archive.org>. The original file is a tar archive containing bziped fragments of text files. The chunks were small, in the order of 1.5MB. In order to ease the load on the cluster filesystem (millions of files

is not a trivial thing) we decided repacked them into bigger chunks, of the order of 100M.

In hindsight this was a bad decision; it would have been adequate if we were the only users of the machine, but on a shared system it would have been better to have more chunks of smaller size, as it would have decreased the granularity of the mapper jobs and made for a more fluid scheduling.

For repacking, gzip compression was used. The bzip2 format is technically superior in that it allows almost-random access to the data, decoupling the chunks from the mapper jobs. However, bzip2 compression is significantly slower than gzip; we considered the time saved when repacking more important than the ability to scale the dataset to a bigger cluster on the spot. In theory, it would have been possible to merge the bzip2 chunks by playing header trickery, but we preferred the safe approach.

For this operation we used the python code under the “extraction” directory. It offers a generator-based API to process the tweets. There is also a script to upload data to a MongoDB instance if needed (we ended up not using MongoDB at all for now).

The repacked chunks were uploaded to the lab cluster HDFS under the /team16 directory, where they were made publically available for other teams. There are also /team16/tweets-small and /team16/tweets-tiny, which are subsets of the main dataset, and can be used for testing purposes.

0.4 Early assessment

Using Java MapReduce jobs proved extremely troublesome. The Hadoop version available was rather obsolete (1.x) and much of the documentation available on the net referenced non-existing mechanisms. Not having administrator access restricted our options; in particular, inclusion of third-party libraries for json parsing was much more difficult than needed. We tried several approaches (fat jars, trickery with the configuration object) to no avail. In the end, it turned out to be much simpler to write jobs using the streaming API. We wrote jobs in Python, as the version installed on the machine already had all the desired libraries.

The streaming API wasn’t free from oddities though. It appears using spaces as end markers for the trigrams wasn’t a great idea, and actually caused the shuffle step to return inconsistent results. Everything went better after we switched to a printable character for the begin/end marker (the dot.)

Every tweet provided a language field, set by the preferences of the user. As a trial of the hadoop system, we counted the languages of our tweets. Results are in table B.

On the contents of every tweet, we computed the simple entropy of the message as a sequence of characters. The entropy per language is presented in figures 1 and 2.

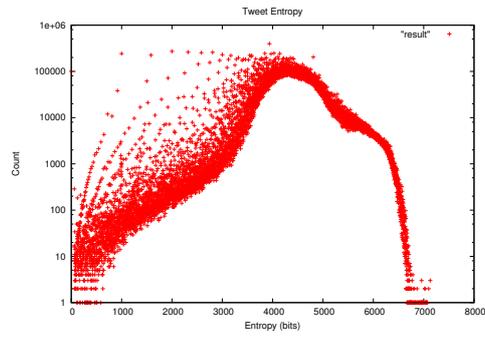


Figure 1: Entropy distribution over tweets

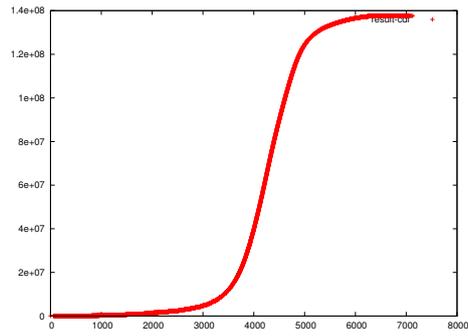


Figure 2: Entropy CDF over tweets

Lang	Count	Lang	Count
ja	2437	bg	2
zh	58	th	1
en	12	pa	1
de	9	he	1
ko	6	ar	1
vi	3		

Table 1: High Entropy Tweets per Language

We suspected entropy might vary per language, and so plotted entropy variation against the language field in ??.

There are a few outliers above 6.6 where we suspected we might find interesting content. The language count for these tweets is reported in ??. It is not surprising that among these outliers are ideogram-based languages, Japanese and Chinese.

Recognizing base64 turned out to be harder than expected. Regex matching yielded a significant number of false positives.

We did, however, identify one odd pattern: sequences of nucleotids (ATGC letters) in long strings. We found a couple of users tweeting these sequences as “Human DNA pieces”. While we initially thought this may be a covert communication channel, the sequences we tested turned out to be genuine human DNA, present in the publicly-available BLAST nucleotid database.

It is straightforward to encode bytes into groups of four nucleotids. With the 140-character twitter limit we can accomodate 35 bytes, or 280 bits, enough for our target bound.

For instance, assigning A,C,G,T the values 0,1,2,3, and writing every ascii value in base 4, lsb first, we can encode `www.facebook.com` as

```
CTCTCTCTCTCTGTGACGCGGACCGATCGCC
CGAGCGTTCGTTCCGGTGTGACGATCGTTCGTC
```

Unfortunately this is not a very good disguise, as such a code would exhibit strong self-correlation of order 2, where for actual genetic data an order 3 would be expected. This could be improved by compressing or encrypting the data with a very simple scheme prior to encoding.

We attempted to match base64-encoded data with the heuristic of attempting a base64 decode, then keeping only the messages for which there is at least one ascii letter. Unfortunately this also produced too many false positives.

0.5 Model building

MapReduce is especially well fitted to compute the probability distributions over a Markov model. We will compute our models both on letters (n-grams) and tokenized words (for the english corpus only). We also need to test several tokenization models. As a part of n-gram model and tokenized words, a MapReduce program has been developed, the logic of the model follows

Mapper (n-gram) Input : The twitter stream, without any pre-processing

Functionality :

1. The twitter stream is parsed and the language field - 'Lang' and Tweet text 'Text' gets extracted.
2. The language fields is used according to the model that we are build [full one or English only]
3. The tweet text further parsed into n-grams, without elimination of any bytes.

Output : Each n-gram will be passed as a key to the reducer. The value is one. n-gram, one

Mapper (tokenized words) :

Input : The twitter stream, without any pre-processing

Functionality :

1. The twitter stream is parsed and the language field - 'Lang' and tweet text 'Text' gets extracted.
2. The language fields is used according to the model that we are build [full one or English only]
3. The tweet text further parsed into words, tweet texts always contains special characters, plural forms which should be eliminated. So regex pattern is used to extract only alpha-numeric characters

Output : Each word will be passed as a key to the reducer. The value is one. word, one

Reducer (common to n-grams and tokens) :

Input : Key, Value

Functionality : Count the Key

Output : Key Count

We hit several trivial (in hindsight) difficulties that required model tweaking, and thus a lot of redundant computations. For the trigrams, the most common word was “RT”: it is used as a marker at the beginning of a tweet to tag retweets. Thus, the model would mostly build strings of the “RT” word repeated. The second most common word was “http” followed closely by “https”. For tokens, we (arbitrarily) excluded hashtags and user handles.

The top 5-grams for the english language are summarized in table 3.

Entropy estimation Once the encoding channel is working, we shall make it run with random input bits, and measure on average how many bits fit in each tweet, for all our models (n-grams, english n-grams, english words). Combined with compression, we will estimate how many messages would be required for typical C&C operations, checking in how many bits we can fit (for instance) a DDoS attack target or a typical shady url.

Detection The difficult part will be extracting suspicious messages conforming too well to our model.

One first technique will be to use different Markov models of different orders, and see if for given accounts, their distribution of tweets follows a low-order model anormally better than a high-order one.

We will also use the n-gram table to compute likelihood for a set of hashtags, and try to extract randomly-generated hashtags from it. We will apply the same n-gram analysis to messages contents, and compare with the first naive regex-based classifier.

Channel building We propose to use an inverted form of Huffman coding to encode arbitrary bits into human-looking chains. The algorithm wil work as follows:

First, prune the model to keep it down to a reasonable size, by excluding uncommon words.

Then, for every state in the chain, build a Huffman tree over the distribution for the next symbol.

To encode, begin with an empty starting state in the Markov chain. Perform a Huffman decoding operation on the tree of the current state, which will consume an arbitrary number of bits and produce a symbol. Output the symbol, compute the new state, and repeat until no bits are left, or the chain picks a terminating symbol, or the maximum output size is exceeded.

To decode, begin with the same empty state; for each token encountered, perform Huffmann compression, repeat until the message is complete.

A Future Work

Another option we want to consider (but less likely to give results) is encoding messages in deliberate grammar or orthography mistakes, through permutation of simple characters. Decoding would be performed by a standard spell checker like aspell.

B Appendix A: results

Lang	Count	Lang	Count
en	46324982	lt	82100
ja	22399870	zh	77609
es	15817824	no	71175
id	10970703	fa	69998
ar	9368831	uk	57837
pt	7189249	he	54388
und	3964165	ur	37947
tr	3897656	is	35487
fr	2947715	hi	19889
ru	2566648	ta	8417
tl	1917014	bn	6140
ko	1730077	ne	5859
th	1664855	hy	2828
it	1029948	ka	2149
nl	1019123	ml	1159
de	725735	te	1071
vi	572220	pa	892
et	498251	bo	746
pl	453255	si	638
sl	429062	my	634
ht	350093	iu	505
sk	263615	gu	492
sv	248050	kn	444
lv	209531	km	433
hu	119744	lo	315
bg	115099	am	196
el	111589	chr	184
fi	110380	dv	146
da	110267	or	47

Table 2: Tweets per language

5gram	count	5gram	count
.the.	7233325	ight.	1943931
.you.	6642536	...m	1943566
...i	4578163	...a	1861019
.and.	4160015	.with	1828758
...i.	3186708	.i.m.	1815360
.for.	2966322	e.you	1750457
...t	2830670	just.	1735798
.that	2691984	.foll	1724977
that.	2569898	with.	1697451
ollow	2305442	...w	1657913
follo	2303823	.like	1655889
...s	2291588	.your	1654147
.just	2044528	.love	1630010
.this	2039137	.have	1627986
this.	2033897		

Table 3: Top 5-grams for english language